

CS5340

Uncertainty Modeling in AI

Lecture 12: Gaussian Processes

“Regression with Bayesian Non-parametrics”

Asst. Prof. Harold Soh

AY 2023/24


Semester 2

Course Schedule (Tentative)

Week	Date	Lecture Topic	Tutorial
1	16 Jan	Introduction to Uncertainty Modeling + Probability Basics	Introduction
2	23 Jan	Simple Probabilistic Models	Introduction and Probability Basics
3	30 Jan	Bayesian networks (Directed graphical models)	More Basic Probability
4	6 Feb	Markov random Fields (Undirected graphical models)	DGM modelling and d-separation
5	13 Feb	Variable elimination and belief propagation	MRF + Sum/Max Product
6	20 Feb	Factor graphs	Quiz 1
-	-	RECESS WEEK	
7	5 Mar	Mixture Models and Expectation Maximization (EM)	Linear Gaussian Models
8	12 Mar	Hidden Markov Models (HMM)	Probabilistic PCA
9	19 Mar	Monte-Carlo Inference (Sampling)	Linear Gaussian Dynamical Systems
10	26 Mar	Variational Inference	MCMC + Langevin Dynamics
11	2 Apr	Inference and Decision-Making (optional)	Diffusion Models + Sequential VAEs
12	9 Apr	Gaussian Processes (optional)	Quiz 2
13	16 Apr	Closing Lecture	Project Presentations


CS5340 in a nutshell

CS5340 is about how to “**represent**” and “**reason**” with **uncertainty** in a computer.



Representation: The *language* is probability and probabilistic graphical models (PGM).

The language is used to **model problems**.



Reasoning: We use learning and inference algorithms to answer questions.

e.g., Belief-propagation/sum-product, MCMC, and variational Bayes

Summary: Sum and Product Rules

- Sum rule:

$$p(x) = \int p(x, y) dy$$

$$p(x) = \sum_y p(x, y)$$

- Product/Chain rule:

$$p(x, y) = p(x|y)p(y)$$

Multivariate Normal Distribution

- Multivariate normal distribution describes a **D -dimensional continuous variable \mathbf{X}** , i.e. $\mathbf{x} \in \mathbb{R}^D$.
- D -dimensional **mean $\boldsymbol{\mu} \in \mathbb{R}^D$** , and $D \times D$ symmetrical positive definite **covariance matrix $\boldsymbol{\Sigma} \in \mathbb{R}_+^{D \times D}$** .

$$p(\mathbf{X} = \mathbf{a} \mid \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{(2\pi)^{D/2} |\boldsymbol{\Sigma}|^{1/2}} \exp\{ -0.5(\mathbf{a} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{a} - \boldsymbol{\mu}) \}, \quad \mathbf{a} \in \mathbb{R}^D$$

Or

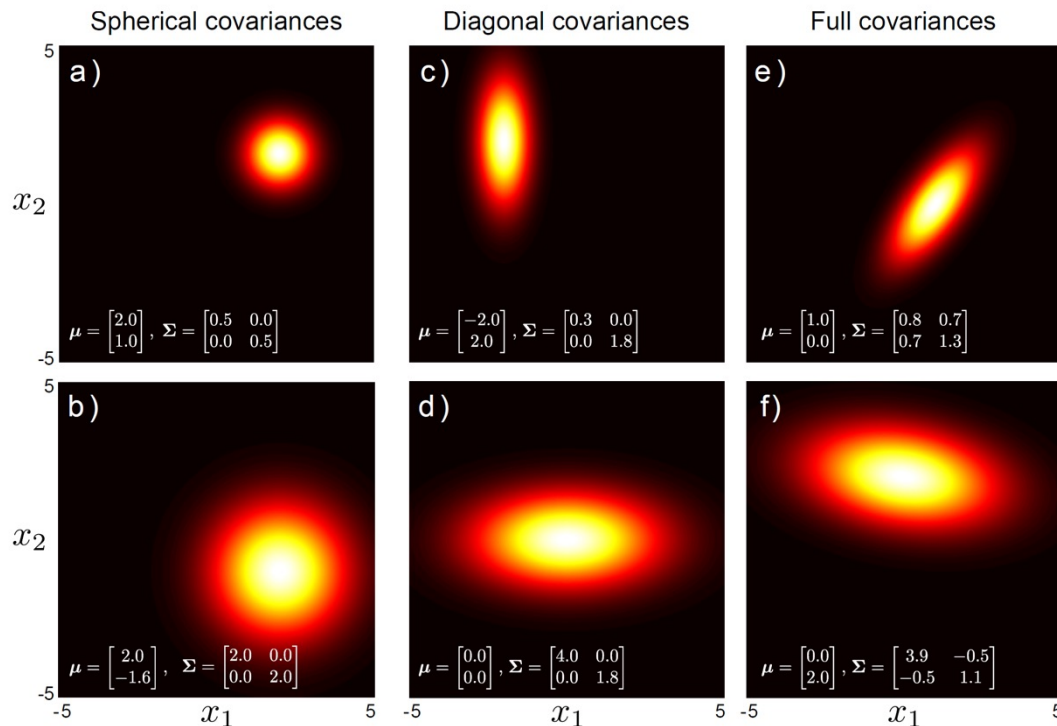
$$p(\mathbf{x}) = \frac{1}{(2\pi)^{D/2} |\boldsymbol{\Sigma}|^{1/2}} \exp\{ -0.5(\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}) \}$$

$$p(\mathbf{x}) = \text{Norm}_{\mathbf{x}}[\boldsymbol{\mu}, \boldsymbol{\Sigma}]$$

Types of Covariance

- Covariance matrix has three forms: **spherical**, **diagonal** and **full**.

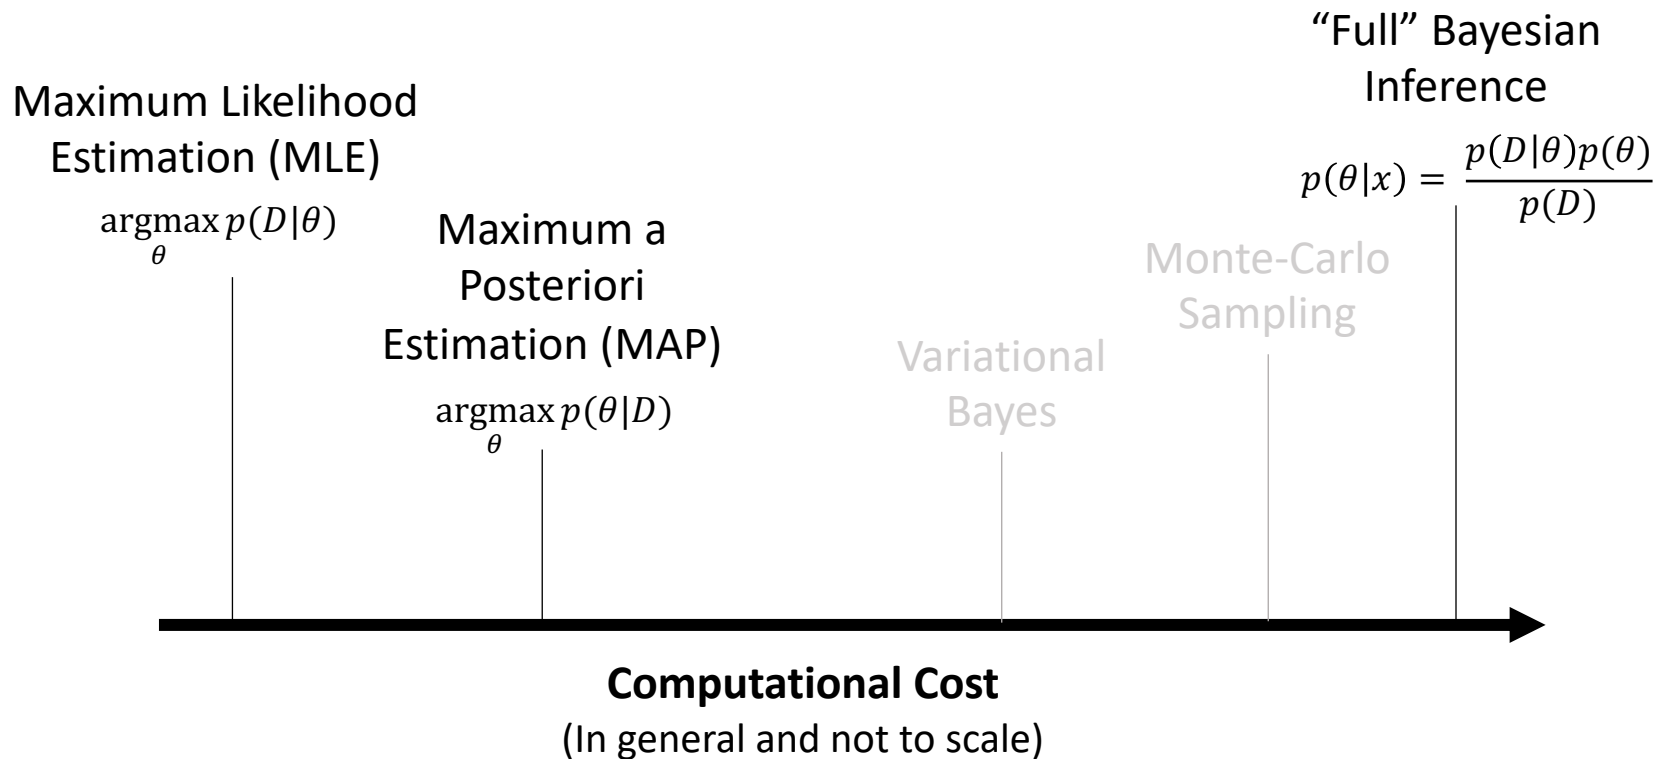
$$\Sigma_{spher} = \begin{bmatrix} \sigma^2 & 0 \\ 0 & \sigma^2 \end{bmatrix} \quad \Sigma_{diag} = \begin{bmatrix} \sigma_1^2 & 0 \\ 0 & \sigma_2^2 \end{bmatrix} \quad \Sigma_{full} = \begin{bmatrix} \sigma_{11}^2 & \sigma_{12}^2 \\ \sigma_{21}^2 & \sigma_{22}^2 \end{bmatrix}$$



Images Source: "Computer Vision: Models, Learning, and Inference", Simon Prince

Learning Parameters

- Common approaches to **learn the unknown parameters** θ from a set of given data $\mathcal{D} = \{x[1], \dots, x[N]\}$:



Learning Outcomes

- Explain the **Gaussian Process** and how it is used for **Regression**
- Describe the different **covariance/kernel functions** and what constitutes a **valid covariance function**.
- Describe the **Empirical Bayesian** procedure for learning **hyperparameters**

Acknowledgements

- Gaussian Processes for Machine Learning, Rasmussen and Williams, 2006. <http://www.gaussianprocess.org>
- Further Exploration:
 - A Visual Exploration of Gaussian Processes, Görtler, et al., Distill, 2019. <https://distill.pub/2019/visual-exploration-gaussian-processes/>
 - Kernel Cookbook, D. Duvenaud, <https://www.cs.toronto.edu/~duvenaud/cookbook/>
 - Gaussian Processes, Neil Lawrence, <http://inverseprobability.com/talks/notes/gaussian-processes.html>

Motivating Example & Applications

Why do we need GPs?


(Non-linear) Regression

- Given $\mathcal{D} = \{(\mathbf{x}[1], y[1]), (\mathbf{x}[2], y[2]), \dots, (\mathbf{x}[N], y[N])\}$
- Want:
 - Function $y = f(\mathbf{x})$
 - Can predict $y^* = f(\mathbf{x}^*)$ for new test point \mathbf{x}^*


Problem: How certain is the prediction $f(\mathbf{x}^*)$?

The Key Idea

- **Bayesian** framework
- **Prior** over possible **functions**
- Infer **Posterior** after seeing data.



(a), prior




(b), posterior

Image Credit: Rasmussen and Williams “Gaussian Processes for Machine Learning”, Chp 2

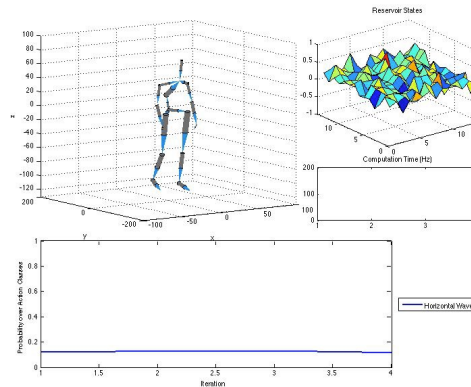
Example Applications

Airline Passenger Predictions




[Wilson and Adams, ICML 2013]. Image from:
<https://people.orie.cornell.edu/andrew/pattern/>

Action Recognition




[Soh and Demiris, 2012]

Object Recognition By touch




[Soh and Demiris, 2015]

Medical Monitoring



[Cheng et al, 2018] Image:
<https://arxiv.org/pdf/1703.09112.pdf>
<https://github.com/bee-hive/MedGP>

Human Trust Modeling




[Soh et al, 2018]

And many more...

GP Code

- Data from Kaggle:
<https://www.kaggle.com/sudalairajkumar/novel-corona-virus-2019-dataset/data#>
- Code on class github



Benefits and Drawbacks of GPs

- **Pros:**

- Conceptually **simple** and **elegant**
- **Interpretable**
- Posterior **Predictive distributions** $p(y^* | \mathbf{x}^*, \mathbf{f})$
- **Flexible**, yet **Prevents Overfitting**
- Model **Selection**

- **Cons:**

- Computationally **expensive**, $O(N^3)$ for basic GP
- Can be **sensitive** to choice of **prior**
 - Covariance/kernel function and hyperparameters can be difficult specify.

The (Wonderful) Properties of Gaussians

Preliminaries

Key ideas

- Gaussians have **nice properties!**
- **“Closed”** under operations of interest
 - Affine transformations, Marginalization, Conditioning
- Probabilistic inference with Gaussians is usually **tractable** and **simple**
 - just linear algebra

Lecture 1: Multivariate Normal Distribution

- Multivariate normal distribution describes a **D -dimensional continuous variable** X , i.e. $\mathbf{x} \in \mathbb{R}^D$.
- D -dimensional **mean** $\boldsymbol{\mu} \in \mathbb{R}^D$, and $D \times D$ symmetrical positive semidefinite **covariance** matrix $\boldsymbol{\Sigma} \in \mathbb{R}_+^{D \times D}$.

$$p(\mathbf{X} = \mathbf{a} \mid \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{(2\pi)^{D/2} |\boldsymbol{\Sigma}|^{1/2}} \exp\{ -0.5(\mathbf{a} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{a} - \boldsymbol{\mu}) \}, \quad \mathbf{a} \in \mathbb{R}^D$$

Or


$$p(\mathbf{x}) = \frac{1}{(2\pi)^{D/2} |\boldsymbol{\Sigma}|^{1/2}} \exp\{ -0.5(\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}) \}$$

$$p(\mathbf{x}) = \text{Norm}_{\mathbf{x}}[\boldsymbol{\mu}, \boldsymbol{\Sigma}]$$

Lecture 1: Types of Covariance

- Covariance matrix has three forms: **spherical**, **diagonal** and **full**.

$$\Sigma_{spher} = \begin{bmatrix} \sigma^2 & 0 \\ 0 & \sigma^2 \end{bmatrix} \quad \Sigma_{diag} = \begin{bmatrix} \sigma_1^2 & 0 \\ 0 & \sigma_2^2 \end{bmatrix} \quad \Sigma_{full} = \begin{bmatrix} \sigma_{11}^2 & \sigma_{12}^2 \\ \sigma_{21}^2 & \sigma_{22}^2 \end{bmatrix}$$



Images Source: "Computer Vision: Models, Learning, and Inference", Simon Prince

Lecture 1: Sum and Product Rules

- Sum rule:

$$p(x) = \int p(x, y) dy$$

$$p(x) = \sum_y p(x, y)$$

- Product/Chain rule:

$$p(x, y) = p(x|y)p(y)$$

The Nice Properties of Gaussians

- Remains “**closed**” under the following operations:
 - Scaling
 - Adding a constant
 - Sum
 - Affine Transformations
 - Marginalization
 - Conditioning
- Applying any of the above operations leads to another **Gaussian**.

Scaling, Adding a Constant, & Sum

If $\mathbf{x} \sim N(\boldsymbol{\mu}, \boldsymbol{\Sigma})$ is a Gaussian random variable then:

- **Scaling:** $\alpha \mathbf{x} \sim N(\alpha \boldsymbol{\mu}, \alpha^2 \boldsymbol{\Sigma})$
- **Adding a Constant:** $\mathbf{x} + \mathbf{a} \sim N(\boldsymbol{\mu} + \mathbf{a}, \boldsymbol{\Sigma})$

Sum: If $\mathbf{x} \sim N(\boldsymbol{\mu}_x, \boldsymbol{\Sigma}_x)$ and $\mathbf{y} \sim N(\boldsymbol{\mu}_y, \boldsymbol{\Sigma}_y)$ are independent then

$$\mathbf{x} + \mathbf{y} \sim N(\boldsymbol{\mu}_x + \boldsymbol{\mu}_y, \boldsymbol{\Sigma}_x + \boldsymbol{\Sigma}_y)$$

Affine Transformation

If $\mathbf{x} \sim N(\boldsymbol{\mu}, \boldsymbol{\Sigma})$ is a Gaussian random variable

and \mathbf{a} is a **constant vector** and \mathbf{B} is a **constant matrix**

then:

$$\mathbf{a} + \mathbf{B}\mathbf{x} \sim N(\mathbf{a} + \mathbf{B}\boldsymbol{\mu}, \mathbf{B}\boldsymbol{\Sigma}\mathbf{B}^T)$$

Marginalization

Let \mathbf{x} and \mathbf{y} be jointly Gaussian random variables.

$$\begin{bmatrix} \mathbf{x} \\ \mathbf{y} \end{bmatrix} \sim N \left(\begin{bmatrix} \boldsymbol{\mu}_x \\ \boldsymbol{\mu}_y \end{bmatrix}, \begin{bmatrix} \mathbf{A} & \mathbf{C} \\ \mathbf{C}^\top & \mathbf{B} \end{bmatrix} \right)$$

Then, the **marginal distribution** of \mathbf{x}

$$\begin{aligned} p(\mathbf{x}) &= \int p(\mathbf{x}, \mathbf{y}) d\mathbf{y} \\ &= N(\boldsymbol{\mu}_x, \mathbf{A}) \end{aligned}$$

Can **simply drop** the **irrelevant variables**.

Exercise: Proof follows from affine property. Consider a \mathbf{B} that “selects” the appropriate variables.

Conditioning

Let \mathbf{x} and \mathbf{y} be jointly Gaussian random variables.

$$\begin{bmatrix} \mathbf{x} \\ \mathbf{y} \end{bmatrix} \sim N \left(\begin{bmatrix} \boldsymbol{\mu}_x \\ \boldsymbol{\mu}_y \end{bmatrix}, \begin{bmatrix} \mathbf{A} & \mathbf{C} \\ \mathbf{C}^\top & \mathbf{B} \end{bmatrix} \right)$$

Then, the **conditional distribution** of \mathbf{x}

$$p(\mathbf{x} | \mathbf{y} = \hat{\mathbf{y}}) =$$

$$N(\boldsymbol{\mu}_x + \mathbf{CB}^{-1}(\hat{\mathbf{y}} - \boldsymbol{\mu}_y), \mathbf{A} - \mathbf{CB}^{-1}\mathbf{C}^\top)$$

The Nice Properties of Gaussians

- Remains “**closed**” under the following operations:
 - Scaling
 - Adding a constant
 - Sum
 - Affine Transformations
 - Marginalization
 - Conditioning
- Applying any of the above operations leads to another **Gaussian**.

Gaussian Processes

Intuition and Formal Definitions

The Key Idea

- **Bayesian** framework
- **Prior** over possible **functions**
- Infer **Posterior** after seeing data.

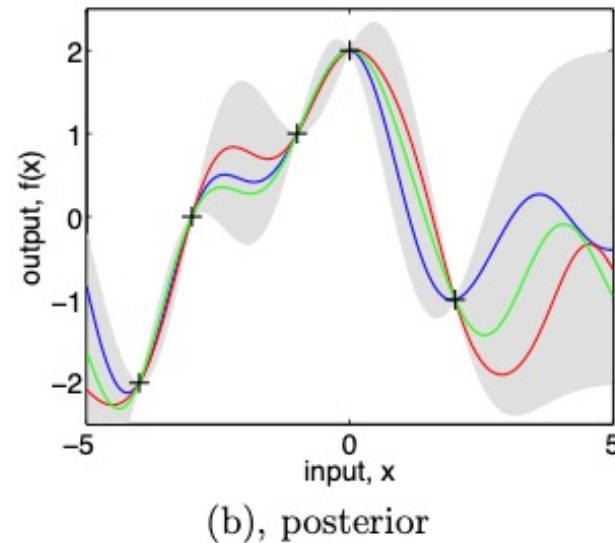



Image Credit: Rasmussen and Williams “Gaussian Processes for Machine Learning”, Chp 2

Intuition for GPs

- Two Approaches / Views
- ➔ • **View 1:** Weight Space View
- **View 2:** Function Space View

(Non-linear) Regression

- Given $\mathcal{D} = \{(\mathbf{x}[1], y[1]), (\mathbf{x}[2], y[2]), \dots, (\mathbf{x}[N], y[N])\}$
- Want:
 - Function $y = f(\mathbf{x})$
 - Can predict $y^* = f(\mathbf{x}^*)$ for new test point \mathbf{x}^*

- Notation:

- **Design Matrix: \mathbf{X}**
- **Targets \mathbf{y}**
- **Data = $\mathcal{D} = (\mathbf{X}, \mathbf{y})$**

$$\mathbf{X} = \begin{bmatrix} | & | & | \\ \mathbf{x}[0] & \mathbf{x}[1] & \mathbf{x}[N] \\ | & | & | \end{bmatrix}$$

$$\mathbf{y} = \begin{bmatrix} y[1] \\ y[2] \\ \vdots \\ y[N] \end{bmatrix}$$

From L3: Linear Regression


- Model for data point with index i :

$$Y[i] = \mathbf{w}^\top \mathbf{x}[i] + \epsilon[i]$$

where:


- $\mathbf{x}[i] = [x[i]_1, x[i]_2, \dots, x[i]_D]^\top$ is a **D-dimensional observed input vector**
- $\mathbf{w} = [w_1, w_2, \dots, w_D]^\top$ is a **coefficient vector**
- $\epsilon[i] \sim N(0, \sigma_n^2)$ is iid zero-mean Gaussian noise

From L3: DGM for Bayesian Linear Regression



- Model uncertainty over \mathbf{w}
- The coefficient vector \mathbf{w} is now a random variable with a prior $p(\mathbf{w}|v) = N(\mathbf{0}, v\mathbf{I})$

From L3: DGM for Bayesian Linear Regression



- Write the **factorization**:

$$p(y[1], \dots, y[N], \mathbf{w}) \\ = p(\mathbf{w}|v) \prod_i p(y[i] | \mathbf{w}^\top \mathbf{x}[i], \sigma_n^2)$$

Exercise: Assume we know σ_n^2 ,
give the MAP solution for \mathbf{w} .

Bayesian Linear Regression

- We want the posterior:

$$p(\mathbf{w}|\mathbf{y}, \mathbf{X}) = \frac{p(\mathbf{y}|\mathbf{X}, \mathbf{w})p(\mathbf{w})}{p(\mathbf{y}|\mathbf{X})}$$

- **Prior:** $p(\mathbf{w}|\Sigma_p) = N(\mathbf{0}, \Sigma_p)$ (Note: consider Σ_p instead of $\nu\mathbf{I}$)

- **Likelihood:** $\prod_i p(y[i]|\mathbf{X}, \mathbf{w}) = \prod_i N(\mathbf{w}^\top \mathbf{x}[i], \sigma_n^2)$

- **Posterior:**

$$p(\mathbf{w}|\mathbf{y}, \mathbf{X}) = N(\mathbf{A}^{-1}\mathbf{X}\mathbf{y}, \mathbf{A}^{-1})$$

$$\text{where } \mathbf{A} = \sigma_n^{-2}\mathbf{X}\mathbf{X}^\top + \Sigma_p^{-1}$$

Bayesian Linear Regression

- **Posterior:**


$$p(\mathbf{w}|\mathbf{y}, \mathbf{X}) = N(\mathbf{A}^{-1}\mathbf{X}\mathbf{y}, \mathbf{A}^{-1})$$

$$\text{where } \mathbf{A} = \sigma_n^{-2}\mathbf{X}\mathbf{X}^\top + \Sigma_p^{-1}$$

- **Posterior Predictive Distribution:**

$$\begin{aligned} p(y^*|\mathbf{x}^*, \mathbf{X}, \mathbf{y}) &= \int p(y^*|\mathbf{x}^*, \mathbf{w})p(\mathbf{w}|\mathbf{y}, \mathbf{X})d\mathbf{w} \\ &= N(\sigma_n^{-2}\mathbf{x}^{*\top}\mathbf{A}^{-1}\mathbf{X}\mathbf{y}, \mathbf{x}^{*\top}\mathbf{A}^{-1}\mathbf{x}^*) \end{aligned}$$

Visually: Posterior Predictive



From L3: Why Linear Regression?

- Basis function “trick”
- Let $\phi(x)$ be some function that transforms x into another vector of “features”
- E.g.:
 - $\phi(x) = [x, x^2, 1]^T$
 - $\phi(x) = [x^p, x^{p-1}, \dots, x^2, x, 1]^T$
- Then, applying the linear model, we get:
 - $Y[i] = \mathbf{w}^T \phi(\mathbf{x}[i]) + \epsilon[i]$
 - For the examples above, this is *polynomial regression*.
- $\phi(x)$ can be more complex:
 - E.g.: $\phi(x[i]) = h(\mathbf{A}\mathbf{x}[i])$ where h is a nonlinear “activation function” (What is this?)

Bayesian Non-linear Regression

- Let $\Phi = \Phi(\mathbf{X})$
 - apply the basis function to all input points \mathbf{X}
 - Also, $\phi^* = \phi(\mathbf{x}^*)$ (apply to test point)
- Then use the same reasoning to get:

$$p(y^* | \mathbf{x}^*, \mathbf{X}, \mathbf{y}) = N(\sigma_n^{-2} \phi^{*\top} \mathbf{A}^{-1} \Phi \mathbf{y}, \phi^{*\top} \mathbf{A}^{-1} \phi^*)$$

where $\mathbf{A} = \sigma_n^{-2} \Phi \Phi^\top + \Sigma_p^{-1}$

Compare with:

$$p(y^* | \mathbf{x}^*, \mathbf{X}, \mathbf{y}) = N(\sigma_n^{-2} \mathbf{x}^{*\top} \mathbf{A}^{-1} \mathbf{X} \mathbf{y}, \mathbf{x}^{*\top} \mathbf{A}^{-1} \mathbf{x}^*)$$

Towards Kernels

- Define:

- $\boldsymbol{\psi}(\mathbf{x}) = \boldsymbol{\Sigma}_p^{1/2} \boldsymbol{\phi}(\mathbf{x})$
- $k(\mathbf{x}, \mathbf{x}') = \boldsymbol{\psi}(\mathbf{x})^\top \boldsymbol{\psi}(\mathbf{x}')$
- $k^* = k(\mathbf{x}^*, \mathbf{x}^*)$

Apply “kernel trick”:
compute inner product
 $k(\mathbf{x}, \mathbf{x}') = \boldsymbol{\psi}(\mathbf{x})^\top \boldsymbol{\psi}(\mathbf{x}')$ by
evaluating a kernel function

$$\mathbf{K} = \begin{bmatrix} k(\mathbf{x}[1], \mathbf{x}[1]) & \cdots & k(\mathbf{x}[1], \mathbf{x}[N]) \\ k(\mathbf{x}[2], \mathbf{x}[1]) & \cdots & k(\mathbf{x}[2], \mathbf{x}[N]) \\ \vdots & \ddots & \vdots \\ k(\mathbf{x}[N], \mathbf{x}[1]) & \cdots & k(\mathbf{x}[N], \mathbf{x}[N]) \end{bmatrix} \quad \mathbf{k}^* = \begin{bmatrix} k(\mathbf{x}[1], \mathbf{x}^*) \\ k(\mathbf{x}[2], \mathbf{x}^*) \\ \vdots \\ k(\mathbf{x}[N], \mathbf{x}^*) \end{bmatrix}$$

Towards Kernels

Rewrite the posterior predictive:

$$\begin{aligned} p(y^* | \mathbf{x}^*, \mathbf{X}, \mathbf{y}) &= N(\sigma_n^{-2} \boldsymbol{\phi}^{*\top} \mathbf{A}^{-1} \boldsymbol{\Phi} \mathbf{y}, \boldsymbol{\phi}^{*\top} \mathbf{A}^{-1} \boldsymbol{\phi}^*) \\ &= N(\boldsymbol{\mu}, \mathbf{V}) \end{aligned}$$


Where:

$$\begin{aligned} \boldsymbol{\mu} &= \mathbf{k}^{*\top} (\mathbf{K} + \sigma_n^2 I)^{-1} \mathbf{y} \\ \mathbf{V} &= k^* - \mathbf{k}^{*\top} (\mathbf{K} + \sigma_n^2 I)^{-1} \mathbf{k}^* \end{aligned}$$

Apply “kernel trick”: compute inner product
 $k(\mathbf{x}, \mathbf{x}') = \boldsymbol{\psi}(\mathbf{x})^\top \boldsymbol{\psi}(\mathbf{x}')$ by evaluating a kernel function

Recap: Weight Space view

Started with **Bayesian Linear Regression**



1. Since all Gaussian, posterior and posterior predictive computable in **closed form**.
2. Using **basis function trick**, perform **non-linear regression**.
3. Using **kernel trick**, obtain **posterior predictive** in terms of **kernel evaluations**

$$p(y^* | \mathbf{x}^*, \mathbf{X}, \mathbf{y}) = N(\boldsymbol{\mu}, \mathbf{V})$$

where:

$$\boldsymbol{\mu} = \mathbf{k}^{*\top} (\mathbf{K} + \sigma_n^2 I)^{-1} \mathbf{y}$$

$$\mathbf{V} = \mathbf{k}^* - \mathbf{k}^{*\top} (\mathbf{K} + \sigma_n^2 I)^{-1} \mathbf{k}^*$$

Intuition for GPs

- Two Approaches / Views
- **View 1:** Weight Space View
- **View 2:** Function Space View



Function space view: Key idea

- Consider the (unknown) function f
- Place a **prior** over f
- And perform **Bayesian inference** to get **posterior** and **posterior predictive**.




Image Credit: Rasmussen and Williams “Gaussian Processes for Machine Learning”, Chp 2

Gaussian Process (GP)

Definition: A **Gaussian process (GP)** is a **collection** of random variables, a **finite number** of which have **joint Gaussian distribution**

We write the GP as

$$f(\mathbf{x}) \sim \mathcal{GP}(m(\mathbf{x}), k(\mathbf{x}, \mathbf{x}'))$$

where

$$\text{Mean function: } m(\mathbf{x}) = \mathbb{E}[f(\mathbf{x})]$$

Covariance function:

$$k(\mathbf{x}, \mathbf{x}') = \mathbb{E}[(f(\mathbf{x}) - m(\mathbf{x}))(f(\mathbf{x}') - m(\mathbf{x}'))]$$


Ingredients of a GP

$$f(x) \sim \mathcal{GP}(m(\mathbf{x}), k(\mathbf{x}, \mathbf{x}'))$$

- **Mean function:** $m(\mathbf{x})$
 - Usually **set** to $m(\mathbf{x}) = \mathbf{0}$
- **Covariance function:** $k(\mathbf{x}, \mathbf{x}')$
 - The **main** ingredient
 - Popular example: **Squared Exponential** (SE) or **Radial Basis Function** (RBF)

Radial Basis Function Kernel

$$k(\mathbf{x}, \mathbf{x}') = \exp\left(-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2l^2}\right)$$

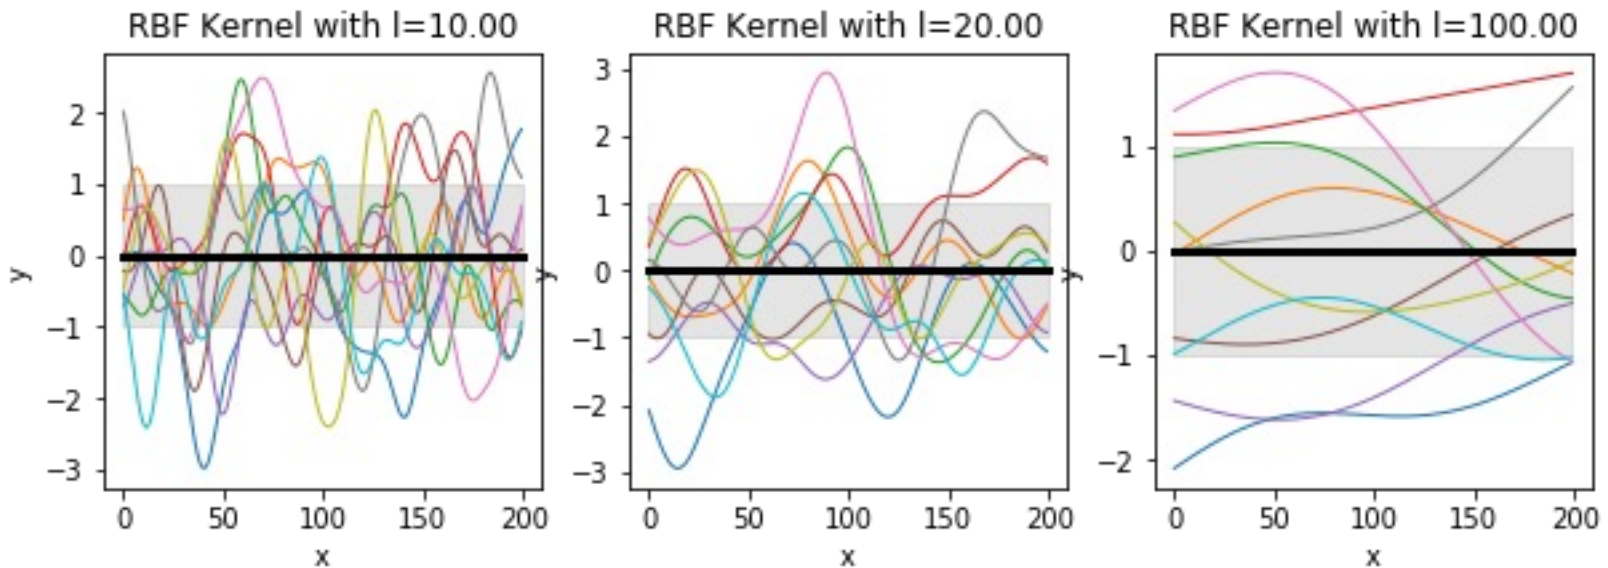


Sampling from the GP prior

- Generate a random vector

$$\mathbf{f}^* \sim N(0, k(\mathbf{X}^*, \mathbf{X}^*))$$

where \mathbf{X}^* is a matrix of input locations



Bayesian Non-linear regression

- We **assume** that:

$$y = f(\mathbf{x}) + \epsilon$$

where $\epsilon \sim N(0, \sigma_n^2)$

- We could obtain a **posterior**:

$$p(f|\mathbf{X}, \mathbf{y}) = \frac{p(\mathbf{y}|\mathbf{X}, f)p(f)}{p(\mathbf{y}|\mathbf{X})}$$

- Often, we really want the **predictive distribution**

$$p(y^*|\mathbf{x}^*, \mathbf{X}, \mathbf{y}) = \int p(y^*|\mathbf{x}^*, f)p(f|\mathbf{y}, \mathbf{X})df$$

Posterior Predictive (noise-free)

The joint distribution of $p(\mathbf{f}, \mathbf{f}^*)$ according to the GP prior is:

$$\begin{bmatrix} \mathbf{f} \\ \mathbf{f}^* \end{bmatrix} \sim N \left(\mathbf{0}, \begin{bmatrix} \mathbf{K} & \mathbf{k}^* \\ \mathbf{k}^{*\top} & k^* \end{bmatrix} \right)$$

Then the **conditional**:

$$p(\mathbf{f}^* | \mathbf{X}^*, \mathbf{X}, \mathbf{f}) = N(\boldsymbol{\mu}, \mathbf{V})$$

where

$$\begin{aligned} \boldsymbol{\mu} &= \mathbf{k}^{*\top} \mathbf{K}^{-1} \mathbf{y} \\ \mathbf{V} &= k^* - \mathbf{k}^{*\top} \mathbf{K}^{-1} \mathbf{k}^* \end{aligned}$$

Posterior Predictive (noisy)

Can only see **noisy** $\mathbf{y} = \mathbf{f} + \epsilon$ where $\epsilon \sim N(\mathbf{0}, \sigma_n^2 \mathbf{I})$

Then, $p(\mathbf{y}, \mathbf{f}^*)$ according to the GP prior is:

$$\begin{bmatrix} \mathbf{y} \\ \mathbf{f}^* \end{bmatrix} \sim N \left(\mathbf{0}, \begin{bmatrix} \mathbf{K} + \sigma_n^2 \mathbf{I} & \mathbf{k}^* \\ \mathbf{k}^{*\top} & k^* \end{bmatrix} \right)$$

Then, the **conditional**:

$$p(\mathbf{f}^* | \mathbf{X}^*, \mathbf{X}, \mathbf{f}) = N(\boldsymbol{\mu}, \mathbf{V})$$

where

$$\boldsymbol{\mu} = \mathbf{k}^{*\top} (\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{y}$$

$$\mathbf{V} = k^* - \mathbf{k}^{*\top} (\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{k}^*$$

From Weight Space View: Towards Kernels

Rewrite the posterior predictive:

$$p(y^* | \mathbf{x}^*, \mathbf{X}, \mathbf{y}) = N(\sigma_n^{-2} \boldsymbol{\phi}^{*\top} \mathbf{A}^{-1} \boldsymbol{\Phi} \mathbf{y}, \boldsymbol{\phi}^{*\top} \mathbf{A}^{-1} \boldsymbol{\phi}^*) \\ = N(\boldsymbol{\mu}, \mathbf{V})$$

Where:

$$\boldsymbol{\mu} = \mathbf{k}^{*\top} (\mathbf{K} + \sigma_n^2 I)^{-1} \mathbf{y} \\ \mathbf{V} = k^* - \mathbf{k}^{*\top} (\mathbf{K} + \sigma_n^2 I)^{-1} \mathbf{k}^*$$

Apply “kernel trick”: compute inner product
 $k(\mathbf{x}, \mathbf{x}') = \boldsymbol{\psi}(\mathbf{x})^\top \boldsymbol{\psi}(\mathbf{x}')$ by evaluating a kernel function

Posterior Predictive (noisy)

$$p(\mathbf{f}^* | \mathbf{X}^*, \mathbf{X}, \mathbf{f}) = N(\boldsymbol{\mu}, \mathbf{V})$$

where

$$\boldsymbol{\mu} = \mathbf{k}^{*\top} (\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{y}$$

$$\mathbf{V} = k^* - \mathbf{k}^{*\top} (\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{k}^*$$

Posterior Predictive: Mean

- “**Linear Predictor**” in terms of the **kernel functions**

$$\mu = \mathbf{k}^{*\top} (\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{y}$$

$$= \mathbf{k}^{*\top} \boldsymbol{\alpha}$$

$$= \sum_{i=1}^N \alpha_i k(\mathbf{x}[i], \mathbf{x}^*)$$

Linear combination of N
kernel functions at the
training points

$$\text{where } \boldsymbol{\alpha} = \underbrace{(\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1}}_{\text{Precision Matrix } \boldsymbol{\Lambda}} \mathbf{y}$$

Precision Matrix $\boldsymbol{\Lambda}$

$\Lambda_{ij} = 0$ iff X_i and X_j are **conditionally independent** given all other X .

Posterior Predictive: Mean

- “**Linear Smoother**” in terms of the **targets**

$$\boldsymbol{\mu} = \mathbf{k}^{*\top} (\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{y}$$

$$= \boldsymbol{\beta}^{*\top} \mathbf{y}$$

$$= \sum_{i=1}^N \beta_i^* y_i$$

Linear combination of
 N targets

$$\text{where } \boldsymbol{\beta}^* = \mathbf{k}^{*\top} (\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1}$$

Posterior Predictive: Variance

$$\mathbf{V} = \underbrace{k^*}_{\text{Prior covariance}} - \underbrace{\mathbf{k}^{*\top} (\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{k}^*}_{\text{Information from other points}}$$

- Does **not** depend on the targets \mathbf{y} .
- This variance is for $p(\mathbf{f}^* | \mathbf{X}^*, \mathbf{X}, \mathbf{f})$
 - Uncertainty about the **true function values**
- To get variance for targets $p(\mathbf{y}^* | \mathbf{X}^*, \mathbf{X}, \mathbf{f})$, add the noise:

$$\mathbf{V}_{\mathbf{y}^*} = \mathbf{V} + \sigma_n^2 \mathbf{I}$$

Computational Complexity

$$p(\mathbf{f}^* | \mathbf{X}^*, \mathbf{X}, \mathbf{f}) = N(\boldsymbol{\mu}, \mathbf{V})$$

where

$$\boldsymbol{\mu} = \mathbf{k}^{*\top} (\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{y}$$

$$\mathbf{V} = k^* - \mathbf{k}^{*\top} (\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{k}^*$$

Complexity is related to cost of inverting $(\mathbf{K} + \sigma_n^2 \mathbf{I})$: $O(N^3)$

Also maintain kernel matrix of size $O(N^2)$

Bayesian Non-parametrics


- GPs are a specific example of **Bayesian non-parametric** models.
- Does **not mean no parameters or no assumptions**
- Number of parameters **grows** with data
 - Theoretically “infinite”
- Can represent **very complex functions**, but Bayesian so, naturally **can prevent overfitting**
- **See also:** Dirichlet Processes

Covariance Functions

*Kernels, Stationary and non-stationary kernels,
Positive Semi-definite kernels*


Ingredients of a GP

$$f(x) \sim \mathcal{GP}(m(\mathbf{x}), k(\mathbf{x}, \mathbf{x}'))$$

- **Mean function:** $m(\mathbf{x})$
 - Usually **set** to $m(\mathbf{x}) = \mathbf{0}$
-  • **Covariance function:** $k(\mathbf{x}, \mathbf{x}')$
 - The **main** ingredient
 - Popular example: **Squared Exponential** (SE) or **Radial Basis Function** (RBF)

Key Ideas


- Covariance functions or kernels encode **similarity**
- Examine some **properties** and **examples**
 - (Non)Stationary, (An)isotropic
- **Not all** similarity functions are **valid** covariance functions
 - Need to be Positive Semi-Definite (PSD)
- Making **new kernels**



Radial Basis Function Kernel


$$k(\mathbf{x}, \mathbf{x}') = \exp\left(-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2l^2}\right)$$

l is a **hyperparameter**



Stationary Kernels

- A **stationary kernel** is a function of $\mathbf{x} - \mathbf{x}'$
 - **Invariant to translations** in input space
 - Example: Matern, Rational Quadratic, Exponential
- **Isotropic** if function of only $\|\mathbf{x} - \mathbf{x}'\|$
 - Invariant to **all rigid motions**.
 - E.g., rotation, reflection
 - Example: RBF



Anisotropic Stationary Kernel

Consider:

$$k(\mathbf{x}, \mathbf{x}') = \exp(-(\mathbf{x} - \mathbf{x}')^\top \mathbf{M}(\mathbf{x} - \mathbf{x}'))$$

where \mathbf{M} is a PSD matrix

Examples:

- $\mathbf{M} = \mathbf{\Psi}$ where $\mathbf{\Psi} = \text{diag} \left([l_j^{-2}]_j^D \right)$:
 - Can specify lengthscale (feature importance) for each dimension.
 - With optimization: Automatic Relevance Determination (ARD)
- $\mathbf{M} = \mathbf{B}\mathbf{B}^\top$ where \mathbf{B} is a $D \times k$ matrix
 - linear **dimensionality reduction**
- $\mathbf{M} = \mathbf{B}\mathbf{B}^\top + \mathbf{\Psi}$
 - **factor analysis** form

Non-Stationary: Dot-Product Kernel

- A **dot-product kernel** is a function of $\mathbf{x}^\top \mathbf{x}'$
 - **Invariant to rotations** (but **not translations**) of input space
 - Obtained by via linear regression

- Can generalize to:

$$k(\mathbf{x}, \mathbf{x}') = \mathbf{x}^\top \mathbf{C} \mathbf{x}'$$

where \mathbf{C} is a general covariance matrix

Positive Semidefinite (PSD) kernels

- The covariance matrix K must be **positive semidefinite**.

$$\mathbf{K} = \begin{bmatrix} k(\mathbf{x}[1], \mathbf{x}[1]) & \cdots & k(\mathbf{x}[1], \mathbf{x}[N]) \\ k(\mathbf{x}[2], \mathbf{x}[1]) & \cdots & k(\mathbf{x}[2], \mathbf{x}[N]) \\ \vdots & \ddots & \vdots \\ k(\mathbf{x}[N], \mathbf{x}[1]) & \cdots & k(\mathbf{x}[N], \mathbf{x}[N]) \end{bmatrix}$$

- A real *symmetric* matrix is **positive semidefinite** iff

$$\forall \mathbf{v} \in \mathbb{R}^d \quad \mathbf{v}^\top \mathbf{K} \mathbf{v} \geq 0$$

- A kernel is PSD if it results in a **PSD matrix**

Is g a valid covariance function?

- A real symmetric matrix is **positive semidefinite** if

$$\forall \mathbf{v} \in \mathbb{R}^d \mathbf{v}^\top \mathbf{K} \mathbf{v} \geq 0$$

*Equivalently, if all eigenvalues of \mathbf{K} are non-negative

- Example:
 - Consider $x, x' \in \mathbb{R}$, is $g(x, x') = x \cdot x'$ a PSD kernel?
 - **Yes** since $\mathbf{v}^\top \mathbf{K} \mathbf{v} = \mathbf{v}^\top \mathbf{X} \mathbf{X}^\top \mathbf{v} = (\mathbf{X}^\top \mathbf{v})^\top (\mathbf{X}^\top \mathbf{v}) \geq 0$

Making New Kernels

- Operations that **preserve PSD**
- **Scaling a kernel** by a non-negative constant $c \geq 0$
 - $k(\mathbf{x}, \mathbf{x}') = ck_1(\mathbf{x}, \mathbf{x}')$
- The **sum of two kernels** is a kernel
 - $k(\mathbf{x}, \mathbf{x}') = k_1(\mathbf{x}, \mathbf{x}') + k_2(\mathbf{x}, \mathbf{x}')$
- The **product of two kernels** is a kernel
 - $k(\mathbf{x}, \mathbf{x}') = k_1(\mathbf{x}, \mathbf{x}')k_2(\mathbf{x}, \mathbf{x}')$
- The **exponentiation of a kernel** is a kernel
 - $k(\mathbf{x}, \mathbf{x}') = \exp(k_1(\mathbf{x}, \mathbf{x}'))$

Can apply these repeatedly, e.g.,

$$k_3(\mathbf{x}, \mathbf{x}') = k_1(\mathbf{x}, \mathbf{x}') + k_2(\mathbf{x}, \mathbf{x}')$$

$$k_5(\mathbf{x}, \mathbf{x}') = k_4(\mathbf{x}, \mathbf{x}') \exp(k_3(\mathbf{x}, \mathbf{x}'))$$

Is g a valid covariance function?

- Can g be constructed from other valid kernels?
- Example:
 - Check: $g(\mathbf{x}, \mathbf{x}') = 2\exp((\mathbf{x}^\top \mathbf{x}')^3 / b) \exp(-\|\mathbf{x} - \mathbf{x}'\|^2)$
where $b > 0$
 - **Yes (Proof as Exercise)**
 - **Hint:** consider the dot-product and RBF kernels. Can we obtain g from using operations on these kernels that preserve PSD?

Hyperparameter Learning

- **Q:** How can we learn the hyperparameters θ ?
- The “right” way:

$$p(f, \theta | \mathbf{X}, \mathbf{y}) = \frac{p(\mathbf{y} | \mathbf{X}, f, \theta) p(f, \theta)}{p(\mathbf{y} | \mathbf{X})}$$

- **Problem:** Intractable in general
- So, approximate by maximizing the log marginal likelihood (“Empirical Bayes”):

$$\log p(\mathbf{y} | \mathbf{X}, \theta)$$

Optimizing the Hyperparameters

The **log marginal likelihood**

$$\log p(\mathbf{y}|\mathbf{X}, \theta) = \log \int p(\mathbf{y}|\mathbf{f}, \mathbf{X}, \theta) p(\mathbf{f}|\theta) d\mathbf{f}$$

Looks difficult, but due to the wonderful properties of Gaussians:

$$p(\mathbf{y}|\mathbf{X}, \theta) = N(\mathbf{0}, \mathbf{K} + \sigma_n^2 \mathbf{I})$$

So,

$$\begin{aligned} \log p(\mathbf{y}|\mathbf{X}, \theta) \\ = -\frac{1}{2} \log \det(\mathbf{K} + \sigma_n^2 \mathbf{I}) - \frac{1}{2} \mathbf{y}^\top (\mathbf{K} + \sigma_n^2 \mathbf{I})^{-1} \mathbf{y} - \frac{N}{2} \log 2\pi \end{aligned}$$

Example: Automatic Relevance Determination

Recall the kernel:

$$k(\mathbf{x}, \mathbf{x}') = \exp(-(\mathbf{x} - \mathbf{x}')^\top \mathbf{M}(\mathbf{x} - \mathbf{x}'))$$

where \mathbf{M} is a PSD matrix

Examples:

- $\mathbf{M} = \mathbf{\Psi}$ where $\mathbf{\Psi} = \text{diag} \left([l_j^{-2}]_j^D \right)$:
 - Can specify lengthscale (feature importance) for each dimension.
 - With optimization: Automatic Relevance Determination (ARD)

Further Exploration of Kernels

- Chapter 4 of “Gaussian Processes for Machine Learning”
 - E.g., kernels for strings

covariance function	expression	S	ND
constant	σ_0^2	✓	
linear	$\sum_{d=1}^D \sigma_d^2 x_d x'_d$		
polynomial	$(\mathbf{x} \cdot \mathbf{x}' + \sigma_0^2)^p$		
squared exponential	$\exp(-\frac{r^2}{2\ell^2})$	✓	✓
Matérn	$\frac{1}{2^{\nu-1}\Gamma(\nu)} \left(\frac{\sqrt{2\nu}}{\ell} r\right)^\nu K_\nu \left(\frac{\sqrt{2\nu}}{\ell} r\right)$	✓	✓
exponential	$\exp(-\frac{r}{\ell})$	✓	✓
γ -exponential	$\exp\left(-\left(\frac{r}{\ell}\right)^\gamma\right)$	✓	✓
rational quadratic	$(1 + \frac{r^2}{2\alpha\ell^2})^{-\alpha}$	✓	✓
neural network	$\sin^{-1} \left(\frac{2\tilde{\mathbf{x}}^\top \Sigma \tilde{\mathbf{x}}'}{\sqrt{(1+2\tilde{\mathbf{x}}^\top \Sigma \tilde{\mathbf{x}})(1+2\tilde{\mathbf{x}}'^\top \Sigma \tilde{\mathbf{x}}')}} \right)$		✓

Image Credit: Rasmussen and Williams “Gaussian Processes for Machine Learning”, Chp 4